

Self Case Study-1_Pump it up-Data Mining the Water Table

Section-I

1. Business Problem

The problem that we are going to solve here is addressing the issue of scarcity of clean drinking water that is being faced by population of Tanzania owing to non-functioning of some of the water-points. Using the given dataset, we have to predict which pumps are functional, which need some repairs, and which don't work at all. We have to predict one of these three classes based on a number of variables about what kind of pump is operating, when it was installed, what is geographic location and how it is managed etc. A smart understanding of which water-points will fail can improve maintenance operations and ensure that clean, potable water is available to communities across Tanzania

2. Mapping the real-world problem to an ML problem

2.1 Description

This particular business problem can be formulated as ML classification problem where the goal is to predict the operating condition of a water-point for each record in the dataset. The target variable/ classes is named as status_group which has three values such as 'functional', 'non-functional' or 'needs repair'

2.2 Data Source

The data can be downloaded from the link: <https://www.drivendata.org/competitions/7/pump-it-up-data-mining-the-water-table/> (<https://www.drivendata.org/competitions/7/pump-it-up-data-mining-the-water-table/>).

We are provided with 4 no. CSV files as follows,

1. train.csv
2. test.csv
3. train_labels.csv
4. SubmissionFormat.csv

The train dataset has 59400 datapoints and 40 features

3 Exploratory Data Analysis

3.1 Importing dependencies

```
In [103]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import re
import time
import warnings
warnings.filterwarnings("ignore")
import numpy as np
from nltk.corpus import stopwords
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.manifold import TSNE
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
import sys
```

3.2 Reading Data

```
In [104]: pd.options.display.max_columns=60 #for reading all columns
```

```
In [105]: df_train = pd.read_csv("train.csv")
```

```
In [106]: df_train.shape
```

```
Out[106]: (59400, 40)
```

```
In [107]: df_train.head() # Loading train dataset
```

```
Out[107]:
```

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name	num_private	basin	subvillage	region	region_c
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.856322	none	0	Lake Nyasa	Mnyusi B	Iringa	
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.147466	Zahanati	0	Lake Victoria	Nyamara	Mara	
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.821329	Kwa Mahundi	0	Pangani	Majengo	Manyara	
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu	0	Ruvuma / Southern Coast	Mahakamani	Mtwara	
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	-1.825359	Shuleni	0	Lake Victoria	Kyanyamisa	Kagera	



```
In [108]: df_train_labels = pd.read_csv("train_lables.csv") #Loading train labels data
```

```
In [109]: df_train_labels.head()
```

```
Out[109]:
```

	id	status_group
0	69572	functional
1	8776	functional
2	34310	functional
3	67743	non functional
4	19728	functional

3.3 Merging df_train & df_train_labels

```
In [110]: df = df_train.merge(df_train_labels, how='left', on='id')
```

3.4 Reading merged data

```
In [111]: print('Number of data points : ', df.shape[0])
print('Number of features : ', df.shape[1])
print('Features : ', df.columns.values)
df.head()
```

Number of data points : 59400

Number of features : 41

Features : ['id' 'amount_tsh' 'date_recorded' 'funder' 'gps_height' 'installer' 'longitude' 'latitude' 'wpt_name' 'num_private' 'basin' 'subvillage' 'region' 'region_code' 'district_code' 'lga' 'ward' 'population' 'public_meeting' 'recorded_by' 'scheme_management' 'scheme_name' 'permit' 'construction_year' 'extraction_type' 'extraction_type_group' 'extraction_type_class' 'management' 'management_group' 'payment' 'payment_type' 'water_quality' 'quality_group' 'quantity' 'quantity_group' 'source' 'source_type' 'source_class' 'waterpoint_type' 'waterpoint_type_group' 'status_group']

Out[111]:

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name	num_private	basin	subvillage	region	region_c
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.856322	none	0	Lake Nyasa	Mnyusi B	Iringa	
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.147466	Zahanati	0	Lake Victoria	Nyamara	Mara	
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.821329	Kwa Mahundi	0	Pangani	Majengo	Manyara	
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu	0	Ruvuma / Southern Coast	Mahakamani	Mtwara	
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	-1.825359	Shuleni	0	Lake Victoria	Kyanyamisa	Kagera	

3.5 Understanding the columns of dataset

```
In [112]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 59400 entries, 0 to 59399
Data columns (total 41 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    59400 non-null  int64
1   amount_tsh           59400 non-null  float64
2   date_recorded        59400 non-null  object
3   funder               55765 non-null  object
4   gps_height           59400 non-null  int64
5   installer            55745 non-null  object
6   longitude            59400 non-null  float64
7   latitude             59400 non-null  float64
8   wpt_name             59400 non-null  object
9   num_private          59400 non-null  int64
10  basin                59400 non-null  object
11  subvillage           59029 non-null  object
12  region               59400 non-null  object
13  region_code          59400 non-null  int64
14  district_code        59400 non-null  int64
15  lga                  59400 non-null  object
16  ward                 59400 non-null  object
17  population           59400 non-null  int64
18  public_meeting       56066 non-null  object
19  recorded_by          59400 non-null  object
20  scheme_management    55523 non-null  object
21  scheme_name          31234 non-null  object
22  permit              56344 non-null  object
23  construction_year    59400 non-null  int64
24  extraction_type       59400 non-null  object
25  extraction_type_group 59400 non-null  object
26  extraction_type_class 59400 non-null  object
27  management           59400 non-null  object
28  management_group     59400 non-null  object
29  payment              59400 non-null  object
30  payment_type         59400 non-null  object
31  water_quality        59400 non-null  object
32  quality_group        59400 non-null  object
33  quantity             59400 non-null  object
34  quantity_group       59400 non-null  object
35  source               59400 non-null  object
36  source_type          59400 non-null  object
37  source_class         59400 non-null  object
38  waterpoint_type      59400 non-null  object
39  waterpoint_type_group 59400 non-null  object
```

```
40 status_group          59400 non-null object  
dtypes: float64(3), int64(7), object(31)  
memory usage: 19.0+ MB
```

```
In [113]: df.isnull().sum()
```

```
Out[113]: id                0
amount_tsh                0
date_recorded             0
funder                   3635
gps_height                0
installer                3655
longitude                 0
latitude                  0
wpt_name                  0
num_private                0
basin                     0
subvillage                371
region                    0
region_code               0
district_code             0
lga                       0
ward                      0
population                0
public_meeting            3334
recorded_by               0
scheme_management         3877
scheme_name               28166
permit                    3056
construction_year         0
extraction_type           0
extraction_type_group     0
extraction_type_class     0
management                0
management_group          0
payment                   0
payment_type              0
water_quality             0
quality_group             0
quantity                  0
quantity_group            0
source                    0
source_type               0
source_class              0
waterpoint_type           0
waterpoint_type_group     0
status_group              0
dtype: int64
```

3. 5.1 Data visualization using library 'pandas_profiling'

```
In [114]: import pandas_profiling as pp
```


In [115]: `pp.ProfileReport(df)`

Summarize dataset: 100%

156/156 [01:41<00:00, 1.35it/s, Completed]

Generate report structure: 100%

1/1 [00:25<00:00, 25.85s/it]

Render HTML: 100%

1/1 [00:12<00:00, 12.08s/it]

Dataset statistics

Number of variables	41
Number of observations	59400
Missing cells	46094
Missing cells (%)	1.9%
Duplicate rows	0
Duplicate rows (%)	0.0%
Total size in memory	19.0 MiB
Average record size in memory	336.0 B

Variable types

Numeric	10
Categorical	29
Boolean	2

Alerts

recorded_by has constant value "GeoData Consultants Ltd"	Constant
date_recorded has a high cardinality: 356 distinct values	High cardinality
funder has a high cardinality: 1897 distinct values	High cardinality
installer has a high cardinality: 2145 distinct values	High cardinality
wpt_name has a high cardinality: 37400 distinct values	High cardinality
subvillage has a high cardinality: 19287 distinct values	High cardinality
lga has a high cardinality: 125 distinct values	High cardinality
ward has a high cardinality: 2092 distinct values	High cardinality
scheme_name has a high cardinality: 2696 distinct values	High cardinality
gps_height is highly correlated with population and 1 other fields (population, construction_year)	High correlation
population is highly correlated with gps_height and 1 other fields (gps_height, construction_year)	High correlation
construction_year is highly correlated with gps_height and 1 other fields (gps_height, population)	High correlation

Out[115]:

3.5.2 Exploring the columns individually one by one for more clarity

3.5.2.1 Column 'amount_tsh'

```
In [116]: df['amount_tsh'].value_counts()
```

```
Out[116]: 0.00      41639
          500.00    3102
          50.00    2472
          1000.00   1488
          20.00    1463
          200.00   1220
          100.00    816
          10.00    806
          30.00    743
          2000.00   704
          250.00   569
          300.00   557
          5000.00   450
          5.00     376
          25.00    356
          3000.00   334
          1200.00   267
          1500.00   197
          6.00     190
          200.00   176
```

```
In [117]: df.loc[df['amount_tsh']==0].groupby('status_group').size()
```

```
Out[117]: status_group
          functional      19706
          functional needs repair  3048
          non functional    18885
          dtype: int64
```

```
In [119]: df.groupby(['amount_tsh', 'status_group']).size().head(20)
```

```
Out[119]: amount_tsh  status_group
0.00      functional      19706
          functional needs repair    3048
          non functional      18885
0.20      non functional         3
0.25      functional          1
1.00      non functional         3
2.00      functional          13
5.00      functional        330
          non functional         46
6.00      functional        174
          functional needs repair     3
          non functional         13
7.00      functional         54
          non functional         15
9.00      non functional         1
10.00     functional        623
          functional needs repair    16
          non functional        167
```

From the above analysis of 'amount_tsh' it is observed that the column has total 41639 zero values which is about 70% of total datapoints. When the static head is zero, the suction and discharge points are at same level. this is in favour of the pump. we have total 19706 functional water points when static head is zero. However there are 18885 non functional water points at zero head.

3.5.2.2 Column 'date_recorded'

```
In [120]: df['date_recorded'].isna().sum()
```

```
Out[120]: 0
```

```
In [121]: df.date_recorded.nunique()
```

```
Out[121]: 356
```

This column has zero null values and total 356 unique values. For now let us keep this column as is.

3.5.2.3 Column 'funder'

```
In [122]: df.funder.nunique()
```

```
Out[122]: 1897
```

```
In [123]: df['funder'].isna().sum()
```

```
Out[123]: 3635
```

this column has 1897 unique values and 3635 missing values

```
In [124]: pd.set_option('display.max_rows', None)
df['funder'].value_counts().head(150).sum()
```

```
Out[124]: 47206
```

```
In [125]: df['funder'].fillna(value='Undefined',inplace=True)
df['funder'].replace(to_replace = '0', value = 'Undefined' , inplace=True) #replacing '0' & missing values with 'Undefined'
```

```
In [127]: df['funder'].value_counts().head(30)
```

```
Out[127]: Government Of Tanzania    9084
          Undefined                 4412
          Danida                   3114
          Hesawa                   2202
          Rwssp                    1374
          World Bank               1349
          Kkkt                    1287
          World Vision             1246
          Unicef                   1057
          Tasaf                    877
          District Council         843
          Dhv                      829
          Private Individual       826
          Dwsp                     811
          Norad                    765
          Germany Republi         610
          Tcrs                     602
          Ministry Of Water       590
          Water                    583
          ...
```

```
In [129]: top_30_funders = ['Government Of Tanzania', 'Undefined', 'Danida', 'Hesawa', 'Rwssp', 'World Bank', 'Kkkt', 'World Vision', 'Unicef', 'Tas
```

```
In [132]: df.loc[~df["funder"].isin(top_30_funders), "funder"] = 'other'
```

```
In [134]: df['funder'].unique()
```

```
Out[134]: array(['other', 'Unicef', 'Dwsp', 'Rwssp', 'Wateraid', 'Danida',
                  'World Vision', 'Hesawa', 'Isf', 'Government Of Tanzania', 'Water',
                  'Private Individual', 'Undefined', 'Lga', 'District Council',
                  'Kkkt', 'Norad', 'Dwe', 'Rc Church', 'Tcrs', 'Germany Republi',
                  'Netherlands', 'Tasaf', 'World Bank', 'Fini Water', 'Dhv', 'Amref',
                  'Ministry Of Water', 'Adb', 'Oxfam', 'Hifab'], dtype=object)
```


```
In [135]: df['funder'].isna().sum()
```

```
Out[135]: 0
```

```
In [136]: df['funder'].value_counts().head(20)
```

```
Out[136]: other 22498
Government Of Tanzania 9084
Undefined 4412
Danida 3114
Hesawa 2202
Rwssp 1374
World Bank 1349
Kkkt 1287
World Vision 1246
Unicef 1057
Tasaf 877
District Council 843
Dhv 829
Private Individual 826
Dwsp 811
Norad 765
Germany Republi 610
Tcrs 602
Ministry Of Water 590
Water 583
Name: funder, dtype: int64
```

```
In [137]: plt.figure(figsize=(28,12))
ax = sns.countplot(x='funder', hue='status_group', data=df)
ax.tick_params(axis='x', rotation=45)
```



As can be seen from the above plot the most of the waterpoints funded by government of tanzania are non-functional

3.5.2.4 Column 'gps_height'

```
In [138]: df['gps_height'].value_counts().head(20)
```

```
Out[138]: 0      20438
-15      60
-16      55
-13      55
1290     52
-20      52
-14      51
303      51
-18      49
-19      47
1295     46
1269     46
1304     45
-23      45
280      44
1538     44
-8       44
1286     44
-17      44
320      43
Name: gps_height, dtype: int64
```


```
In [139]: df['gps_height'].isna().sum()
```

```
Out[139]: 0
```



```
In [140]: plt.figure(figsize=(28,12))  
sns.kdeplot(data=df, x='gps_height', hue="status_group", gridsize=200)
```

```
Out[140]: <AxesSubplot:xlabel='gps_height', ylabel='Density'>
```



3.5.2.5 Column 'installer'

```
In [141]: df['installer'].value_counts().head(100).sum()
```

```
Out[141]: 43663
```

```
In [142]: df['installer'].nunique()
```

```
Out[142]: 2145
```

```
In [143]: df['installer'].isna().sum()
```

```
Out[143]: 3655
```

```
In [144]: df['installer'].fillna(value='Undefined',inplace=True)  
df['installer'].replace(to_replace = '0', value ='Undefined' , inplace=True) #replacing '0' category with 'Undefined'
```

```
In [145]: pd.set_option('display.max_rows', None)  
df['installer'].value_counts().head(20)
```

```
Out[145]: DWE                17402  
Undefined            4432  
Government          1825  
RWE                 1206  
Commu               1060  
DANIDA              1050  
KKKT                 898  
Hesawa              840  
TCRS                707  
Central government   622  
CES                 610  
Community           553  
DANID               552  
District Council    551  
HESAWA              539  
LGA                 408  
World vision        408  
WEDECO              397  
TASAF               396  
...
```

```
In [146]: df['installer'].replace(to_replace = ("Gove","Gover","GOVERN", "GOVERN", "GOVERNME", "Governmen","Government","GOVER"), value ="G
df['installer'].replace(to_replace = ("RW","RWE","RWE /Community","RWE Community","RWE/ Community","RWE/Community","RWE/DWE","RWE
df['installer'].replace(to_replace = ("Commu", "Communit", "Community", "COMMUNITY BANK","Comunity"), value ="Community", inplace
df['installer'].replace(to_replace = ("Danda","DANIAD","Danid","DANIDA","DANIDA CO","DANIDS","DANNIDA","DANID"), value ="DANIDA",
df['installer'].replace(to_replace = ("Cebtral Government","Cental Government","Centr","Centra Government","Centra govt","Central
df['installer'].replace(to_replace = ("COUN","Counc","Council","Distri","District Council","District Community j","District Coun
df['installer'].replace(to_replace = ("Hesawa","HESAW","Hesewa","HESAWA"),value ='HESAWA' , inplace=True)
df['installer'].replace(to_replace = ("World Division","World Visiin","World vision","World Vission","World Vision"),value ='World
df['installer'].replace(to_replace = ("Distric Water Department","District Water Department","District water depar","District wat
df['installer'].replace(to_replace = ("FINN WATER","FinW","FinWate","FinWater","Fini water","Fini Water" ),value ='Fini Water' ,
df['installer'].replace(to_replace = ("RC","RC .Church","RC C","RC Ch","RC Churc","RC Church","RC CHURCH BROTHER","RC church/CEFA
df['installer'].replace(to_replace = ("Villa","VILLAGER","Villagerd","Villagers","Villages","Village Council","Villi","villigers")
```

```
In [147]: df['installer'].value_counts().head(30)
```

```
Out[147]: DWE                17402
Undefined            4432
Government          2592
DANIDA              1679
Community           1670
HESAWA              1381
RWE                 1306
District Council    1162
Central Government  1080
KKKT                898
TCRS                707
World Vision        693
CES                 610
Fini Water          553
RC Church           490
LGA                 408
WEDECO              397
TASAF               396
AMREF               329
TURCA              312
```

```
In [148]: top_30_installer = ["DWE","Undefined","Government","DANIDA","Community","HESAWA","RWE","District Council","Central Government","K
```

```
In [149]: df.loc[~df["installer"].isin(top_30_installer), "installer"] = 'other'
```


```
In [150]: df['installer'].nunique()
```

```
Out[150]: 31
```

```
In [151]: df['installer'].isna().sum()
```

```
Out[151]: 0
```

```
In [152]: plt.figure(figsize=(28,12))  
ax = sns.countplot(x='installer', hue='status_group', data=df)  
ax.tick_params(axis='x', rotation=45)
```



From the above plot it can be seen that most of the waterpoints install by government are non-functional.

3.5.2.6 Column 'longitude & latitude'

```
In [153]: df['longitude'].isna().sum()
```

```
Out[153]: 0
```

```
In [154]: df['latitude'].isna().sum()
```

```
Out[154]: 0
```

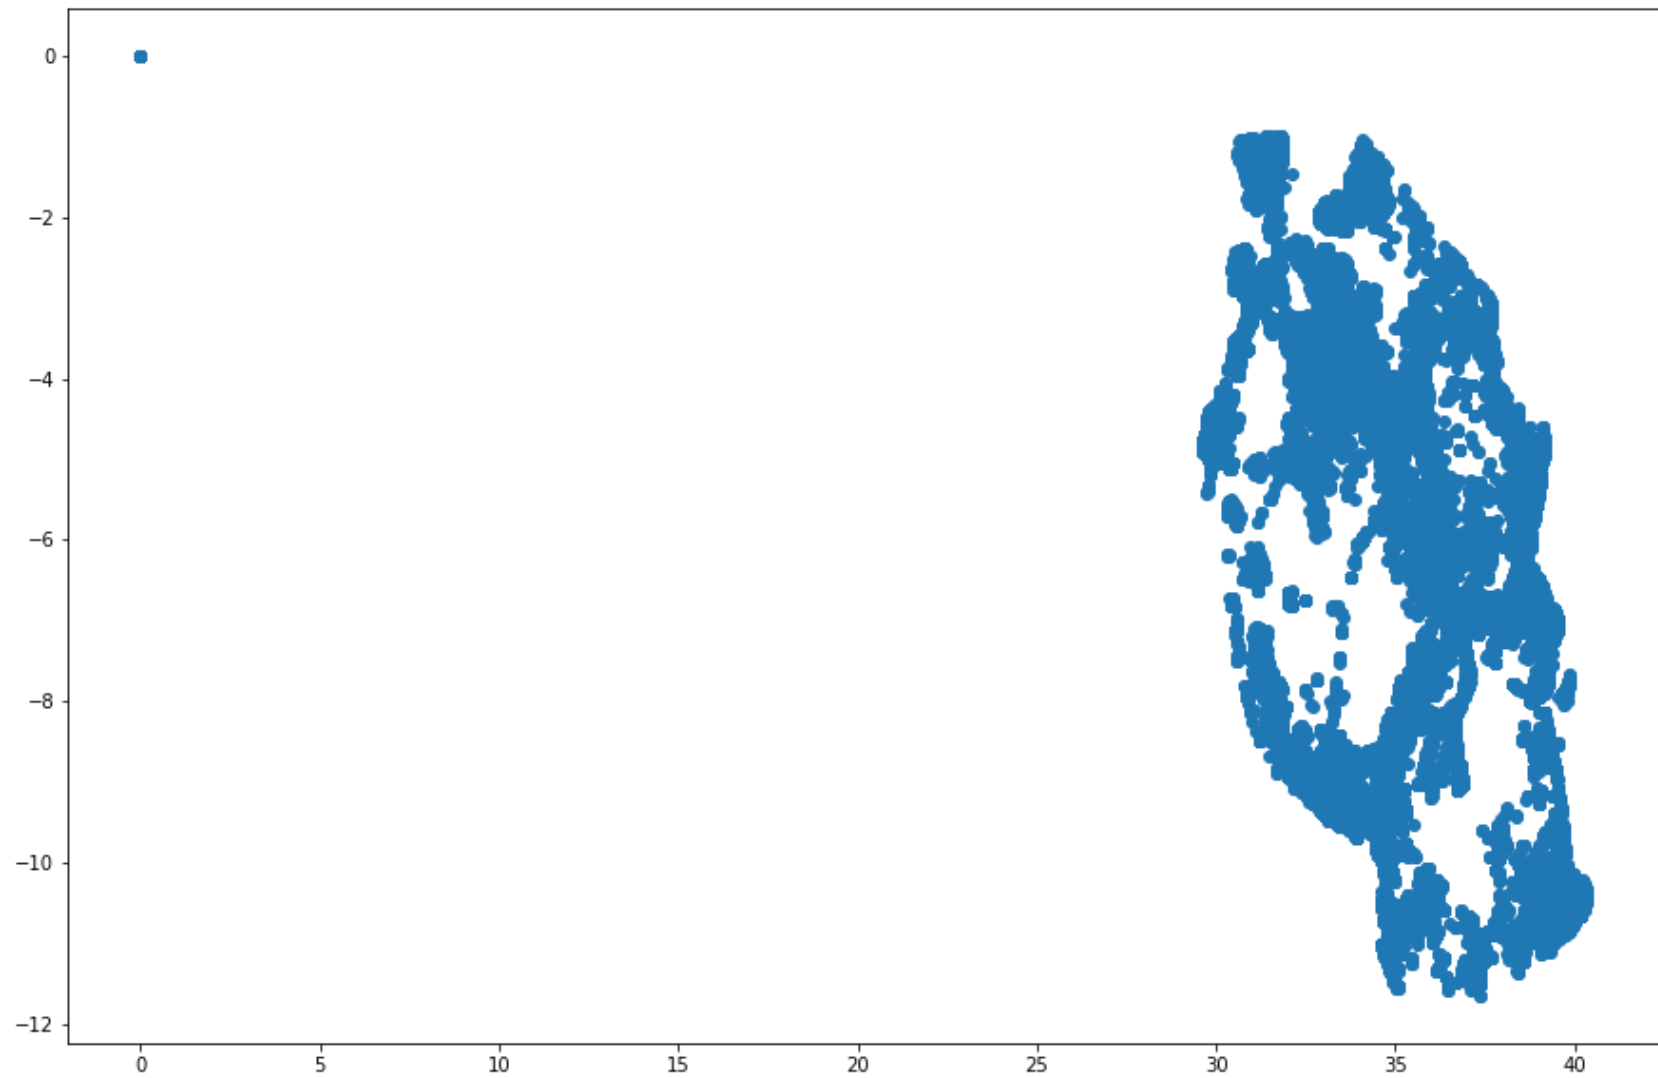
```
In [155]: df['longitude'].nunique()
```

```
Out[155]: 57516
```

```
In [156]: df['latitude'].nunique()
```

```
Out[156]: 57517
```

```
In [157]: plt.figure(figsize=(15,10))
plt.scatter(x="longitude", y="latitude", data=df)
plt.show()
```




From the above scatterplot it is seen that there is outlier at 0. we will replace the same with 'mean'

```
In [158]: df['longitude'].mean()
```

```
Out[158]: 34.07742669202832
```

```
In [159]: df['longitude'].replace(to_replace = 0 , value =34.07742669202832 , inplace=True)
```

```
In [160]: plt.figure(figsize=(20,10))  
plt.scatter(x="longitude", y="latitude", data=df)  
plt.show()
```



3.5.2.7 Column 'wpt_name'

```
In [161]: df['wpt_name'].isna().sum()
```

```
Out[161]: 0
```

```
In [162]: df['wpt_name'].nunique()
```

```
Out[162]: 37400
```

```
In [164]: df['wpt_name'].value_counts().head(20)
```

```
Out[164]: none                3563
Shuleni                    1748
Zahanati                   830
Msikitini                  535
Kanisani                   323
Bombani                    271
Sokoni                     260
Ofisini                    254
School                    208
Shule Ya Msingi           199
Shule                      152
Sekondari                  146
Muungano                   133
Mkombozi                   111
Madukani                   104
Mbugani                    94
Hospital                   94
Upendo                     93
Kituo Cha Afya            90
Mkuyuni                    88
Name: wpt_name, dtype: int64
```


```
In [165]: top_20_wpt = ["none", "Shuleni", "Zahanati", "Msikitini", "Kanisani", "Bombani", "Sokoni", "Ofisini", "School", "Shule Ya Msingi", "Shule",
```

```
In [170]: w = list(df.loc[~df["wpt_name"].isin(top_20_wpt), "wpt_name"])
df['wpt_name'].replace(to_replace = w , value = 'other' , inplace=True)
```

```
In [171]: df["wpt_name"].nunique()
```

```
Out[171]: 21
```

```
In [172]: plt.figure(figsize=(28,12))
ax = sns.countplot(x='wpt_name', hue="status_group", data=df)
ax.tick_params(axis='x', rotation=45)
```



"wpt_name" has 37400 unique values so it does not make sense to retain this feature. the same need be dropped.

```
In [173]: df.drop(columns='wpt_name', inplace=True)
```

3.5.2.8 Column 'num_private'

```
In [174]: df['num_private'].nunique()
```

```
Out[174]: 65
```

```
In [175]: df['num_private'].value_counts()
```

```
Out[175]: 0      58643
          6       81
          1       73
          5       46
          8       46
          32      40
          45      36
          15      35
          39      30
          93      28
           3      27
           7      26
           2      23
          65      22
          47      21
           4      20
         102      20
          17      17
          80      15
          ..      ..
```

most of the values in "num_private" are zero. hence we will remove this features

```
In [176]: df.drop(columns='num_private',inplace=True )
```

3.5.2.9 Column 'basin'

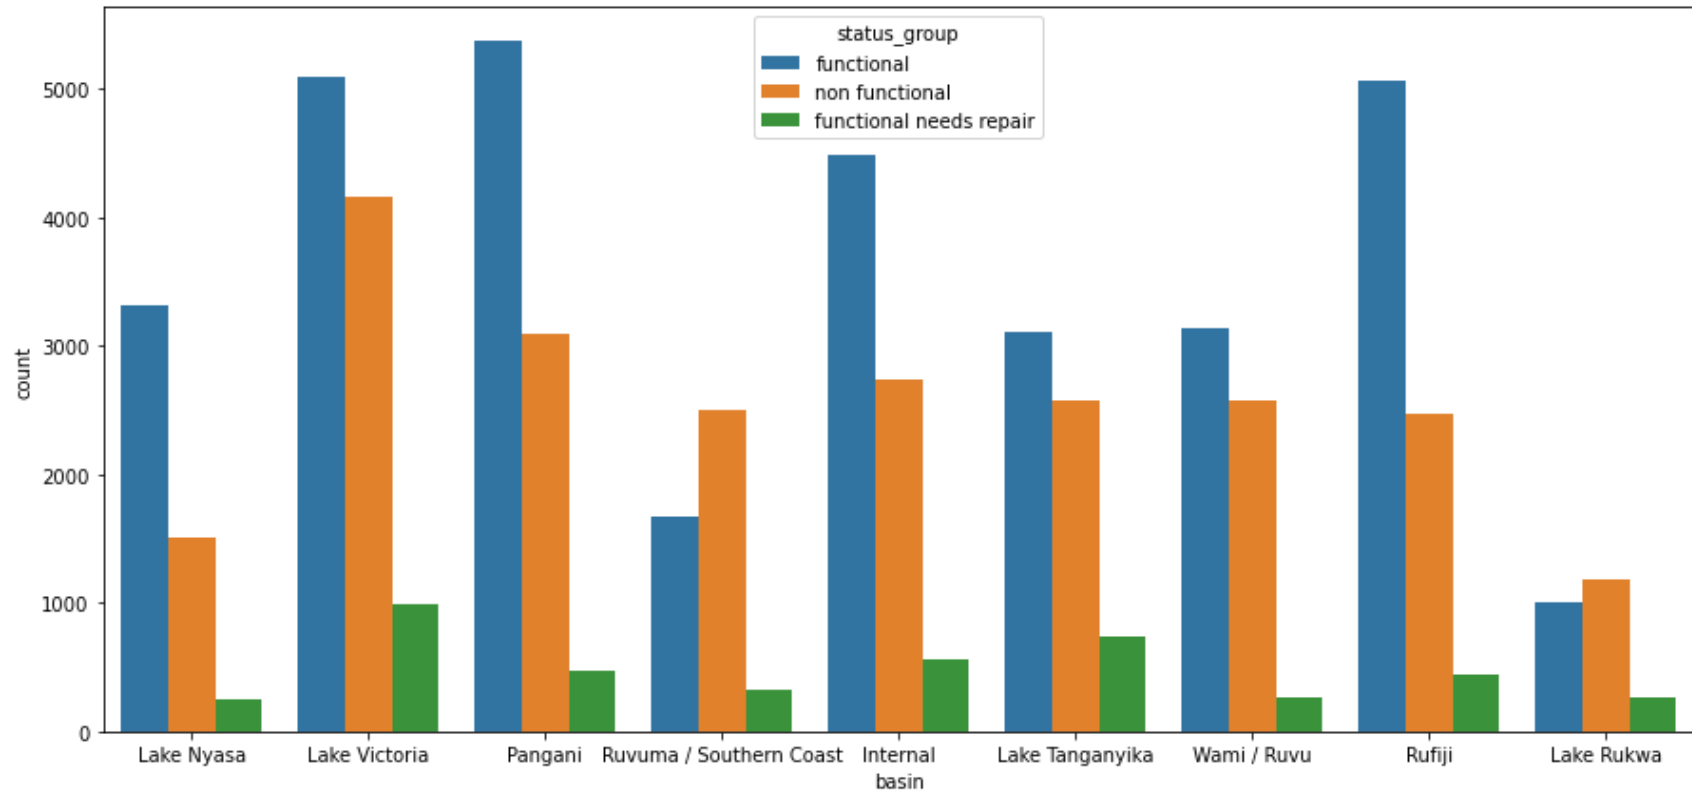
```
In [177]: df['basin'].value_counts()
```

```
Out[177]: Lake Victoria      10248
          Pangani             8940
          Rufiji              7976
          Internal            7785
          Lake Tanganyika     6432
          Wami / Ruvu         5987
          Lake Nyasa          5085
          Ruvuma / Southern Coast 4493
          Lake Rukwa          2454
          Name: basin, dtype: int64
```

```
In [178]: df['basin'].isna().sum()
```

```
Out[178]: 0
```

```
In [179]: plt.figure(figsize=(15,7))  
ax = sns.countplot(x='basin', hue="status_group", data=df)
```



This feature seems have good corelation with class variable 'status_group'

3.5.2.10 Column 'subvillage', 'region'

```
In [180]: df['subvillage'].isna().sum()
```

```
Out[180]: 371
```

```
In [181]: df['subvillage'].nunique()
```

```
Out[181]: 19287
```

```
In [182]: df['subvillage'].value_counts().head(30)
```

```
Out[182]: Madukani      508
          Shuleni       506
          Majengo       502
          Kati          373
          Mtakuja       262
          Sokoni        232
          M             187
          Muungano      172
          Mbuyuni       164
          Mlimani       152
          Songambele    147
          Msikitini     134
          Miembeni      134
          1             132
          Kibaoni       114
          Kanisani      111
          I             109
          Mapinduzi     109
          Mjimwema      108
          Mjini         108
          Mkwajuni      104
          Mwenge        102
          Mabatini      98
          Azimio        98
          Mission       95
          Mbugani       95
          Bwawani       91
          Bondeni       90
          Chang'Ombe    88
          Zahanati      86
          Name: subvillage, dtype: int64
```

```
In [183]: df['region'].isna().sum()
```

```
Out[183]: 0
```

```
In [184]: df['region'].nunique()
```

```
Out[184]: 21
```

```
In [185]: df['region'].value_counts()
```

```
Out[185]: Iringa          5294  
Shinyanga          4982  
Mbeya              4639  
Kilimanjaro        4379  
Morogoro            4006  
Arusha              3350  
Kagera              3316  
Mwanza              3102  
Kigoma              2816  
Ruvuma              2640  
Pwani               2635  
Tanga               2547  
Dodoma              2201  
Singida             2093  
Mara                1969  
Tabora              1959  
Rukwa               1808  
Mtwara              1730  
Manyara             1583  
Lindi               1546  
Dar es Salaam        805  
Name: region, dtype: int64
```

```
In [189]: df.groupby(['region', 'subvillage']).size().head(20)
```

```
Out[189]: region  subvillage
Arusha  Afya      15
        Ahara      1
        Alairataat  3
        Alakirikir  4
        Alasai      1
        Aleilelai   4
        Alsini      2
        Ambara      1
        Ambureni    8
        Arahati     4
        Arashi      3
        Arati       5
        Arauyo      21
        Ariahati     1
        Arkaria     5
        Arudeko     11
        Ascarida     1
        Athin Kati   1
        Athni Mwisho 1
        Atsin       1
dtype: int64
```

'subvillage' and 'region' both provide information about location of wells. as 'subvillage' and more number of categorical values we will drop column 'subvillage'

```
In [190]: df.drop(columns='subvillage', inplace=True )
```

```
In [191]: df.shape
```

```
Out[191]: (59400, 38)
```

3.5.2.12 Column 'region_code', 'district_code'

```
In [192]: df['region_code'].nunique()
```

```
Out[192]: 27
```

```
In [193]: df['district_code'].nunique()
```

```
Out[193]: 20
```

```
In [194]: df['region_code'].value_counts()
```

```
Out[194]: 11    5300
          17    5011
          12    4639
           3    4379
           5    4040
          18    3324
          19    3047
           2    3024
          16    2816
          10    2640
           4    2513
           1    2201
          13    2093
          14    1979
          20    1969
          15    1808
           6    1609
          21    1583
          80    1238
          60    1025
          90     917
           7     805
          99     423
           9     390
          24     326
           8     300
          40        1
          Name: region_code, dtype: int64
```




```
In [195]: df['district_code'].value_counts()
```

```
Out[195]: 1      12203
          2      11173
          3       9998
          4       8999
          5       4356
          6       4074
          7       3343
          8       1043
         30        995
         33        874
         53        745
         43        505
         13        391
         23        293
         63        195
         62        109
         60         63
          0         23
         80         12
         67          6
Name: district_code, dtype: int64
```


```
In [196]: df.groupby(['region_code', 'district_code']).size()
```

```
Out[196]: region_code  district_code
1                0           23
              1          888
              3          361
              4          347
              5          358
              6          224
2                1          189
              2         1206
              3          109
              5          201
              6          310
              7         1009
3                1          595
              2          519
              3          877
              4         1225
              5          620
              6          109
              7           12
```

```
In [197]: plt.figure(figsize=(15,7))
ax = sns.countplot(x='district_code', hue="status_group", data=df)
```



```
In [198]: plt.figure(figsize=(15,7))
ax = sns.countplot(x='region_code', hue="status_group", data=df)
```



For now we will keep the feature "district_code" and remove "region_code"

```
In [199]: df.drop(columns='region_code',inplace=True )
```

```
In [200]: df.shape
```

```
Out[200]: (59400, 37)
```

3.5.2.13 Column 'lga', 'ward'

```
In [201]: df['lga'].isna().sum()
```

```
Out[201]: 0
```

```
In [202]: df['ward'].isna().sum()
```

```
Out[202]: 0
```

```
In [203]: df['lga'].value_counts().head(20)
```

```
Out[203]: Njombe          2503
Arusha Rural      1252
Moshi Rural       1251
Bariadi           1177
Rungwe            1106
Kilosa            1094
Kasulu            1047
Mbozi             1034
Meru              1009
Bagamoyo          997
Singida Rural     995
Kilombero         959
Same              877
Kibondo           874
Kyela             859
Kahama            836
Magu              824
Kigoma Rural      824
Maswa             809
Karagwe           771
Name: lga, dtype: int64
```

```
In [204]: df['lga'].nunique()
```

```
Out[204]: 125
```

```
In [205]: df['ward'].nunique()
```

```
Out[205]: 2092
```

```
In [206]: df['ward'].isna().sum()
```

```
Out[206]: 0
```

```
In [207]: df['ward'].value_counts().head(20)
```

```
Out[207]: Igosi          307
Imalinyi          252
Siha Kati         232
Mdandu            231
Nduruma           217
Mishamo           203
Kitunda           203
Msindo            201
Chalinze          196
Maji ya Chai      190
Usuka             187
Ngarenanyuki      172
Chanika           171
Vikindu           162
Mtwango           153
Matola            145
Zinga/Ikerege     141
Wanging'ombe      139
Maramba           139
Itete             137
Name: ward, dtype: int64
```


```
In [208]: df.groupby(["lga", "ward"]).size().head(50)
```

```
Out[208]: lga      ward      size
Arusha Rural Bangata      33
           Bwawani      37
           Ilkiding'a     86
           Kimnyaki      79
           Kiranyi     115
           Kisongo      33
           Mateves      22
           Mlangarini    92
           Moivo       44
           Moshono      44
           Murieti      29
           Musa        29
           Mwandeti     17
           Nduruma     205
           Oldonyosambu   77
           Oljoro        8
           Olkokola     133
           Oltroto      75
           Oltrumet     52
           Sokoni II     42
Arusha Urban Baraa        2
           Daraja Mbili   3
           Elerai       11
           Engutoto       1
           Kaloleni       5
           Kimandolu       2
           Lemara         4
           Levolosi        2
           Ngarenaro       3
           Olorien         4
           Sekei           3
           Sokon I         4
           Sombetini       4
           Terrat         8
           Themi          1
           Unga Ltd        6
Babati      Arri         19
           Bashinet      19
           Bonga         11
           Dabil         47
           Dareda        52
           Duru          15
           Gidas         19
           Madunga       24
```

Magara	38
Magugu	40
Mamire	115
Mwada	15
Nkaiti	22
Qash	13

dtype: int64

```
In [209]: plt.figure(figsize=(15,7))
ax = sns.countplot(x='lga', hue="status_group", data=df)
ax.tick_params(axis='x', rotation=90)
```



For now we will remove 'ward' and keep 'lga'

```
In [210]: df.drop(columns='ward',inplace=True )
```

3.5.2.14 Column 'population'

```
In [211]: df["population"].isna().sum()
```

```
Out[211]: 0
```

```
In [212]: df["population"].nunique()
```

```
Out[212]: 1049
```

```
In [213]: df["population"].value_counts().head(20)
```

```
Out[213]: 0      21381
          1       7025
          200     1940
          150     1892
          250     1681
          300     1476
          100     1146
           50     1139
          500     1009
          350      986
          120      916
          400      775
           60      706
           30      626
           40      552
           80      533
          450      499
           20      462
          600      438
          230      388
Name: population, dtype: int64
```

This feature does not have missing values but most of the values are zero. let us explore more


```
In [214]: df.loc[df["population"]!=0].describe()
```

Out[214]:

	id	amount_tsh	gps_height	longitude	latitude	district_code	population	construction_year
count	38019.000000	38019.000000	38019.000000	38019.000000	38019.000000	38019.000000	38019.000000	38019.000000
mean	37107.559115	447.787681	969.889634	36.074387	-6.139781	6.299456	281.087167	1961.399721
std	21406.803661	3706.770967	612.544787	2.586779	2.737733	11.303334	564.687660	263.994165
min	1.000000	0.000000	-90.000000	29.607122	-11.649440	1.000000	1.000000	0.000000
25%	18514.500000	0.000000	347.000000	34.715340	-8.388839	2.000000	40.000000	1986.000000
50%	37128.000000	0.000000	1135.000000	36.706815	-5.750877	3.000000	150.000000	2000.000000
75%	55505.500000	100.000000	1465.000000	37.940149	-3.597016	5.000000	324.000000	2008.000000
max	74247.000000	350000.000000	2770.000000	40.345193	-1.042375	67.000000	30500.000000	2013.000000

We will replace zeros with mean value

```
In [215]: df['population'].replace(to_replace = 0, value = 281.087167 , inplace=True)
```

```
In [216]: df["population"].describe()
```

Out[216]:

count	59400.000000
mean	281.087167
std	451.765813
min	1.000000
25%	100.000000
50%	281.087167
75%	281.087167
max	30500.000000

Name: population, dtype: float64


```
In [217]: df["population"].value_counts().head(20)
```

```
Out[217]: 281.087167    21381
          1.000000     7025
          200.000000    1940
          150.000000    1892
          250.000000    1681
          300.000000    1476
          100.000000    1146
           50.000000    1139
          500.000000    1009
          350.000000     986
          120.000000     916
          400.000000     775
           60.000000     706
           30.000000     626
           40.000000     552
           80.000000     533
          450.000000     499
           20.000000     462
          600.000000     438
          230.000000     388
Name: population, dtype: int64
```

```
In [389]: #df.groupby(['population', 'status_group']).size()
```

```
In [218]: plt.figure(figsize=(28,12))  
sns.kdeplot(data=df, x='population', hue="status_group", gridsize=200)
```

```
Out[218]: <AxesSubplot:xlabel='population', ylabel='Density'>
```



From the above kd plot it is understood that the more populated area is, higher the chances of the waterpoint being functional

3.5.2.14 Column 'public_meeting'

```
In [219]: df['public_meeting'].value_counts()
```

```
Out[219]: True      51011  
False      5055  
Name: public_meeting, dtype: int64
```

```
In [220]: df['public_meeting'].isna().sum()
```

```
Out[220]: 3334
```

there are 3334 missing values, we will replace those with most occurring values i.e. 'True'

```
In [221]: df['public_meeting'].fillna(value=True, inplace=True)
```

```
In [222]: df['public_meeting'].value_counts()
```

```
Out[222]: True      54345  
False      5055  
Name: public_meeting, dtype: int64
```

3.5.2.14 Column 'recorded_by'

```
In [223]: df['recorded_by'].isna().sum()
```

```
Out[223]: 0
```

```
In [224]: df['recorded_by'].value_counts()
```

```
Out[224]: GeoData Consultants Ltd    59400  
Name: recorded_by, dtype: int64
```

This featured will not be useful as it has only one value. hence we will remove the same.

```
In [225]: df.drop(columns='recorded_by', inplace = True)
```

```
In [226]: df.shape
```

```
Out[226]: (59400, 35)
```

3.5.2.15 Column "scheme_management", "scheme_name", "management", "management_group"

```
In [227]: print(df['scheme_management'].isna().sum())  
print(df['scheme_management'].nunique())  
df['scheme_management'].value_counts()
```

```
3877
```

```
12
```

```
Out[227]: VWC          36793  
WUG          5206  
Water authority  3153  
WUA          2883  
Water Board    2748  
Parastatal     1680  
Private operator 1063  
Company        1061  
Other          766  
SWC            97  
Trust          72  
None           1  
Name: scheme_management, dtype: int64
```

```
In [301]: print(df['scheme_name'].isna().sum())  
print(df['scheme_name'].nunique())  
df['scheme_name'].value_counts()
```

```
In [229]: print(df['management'].isna().sum())
print(df['management'].nunique())
df['management'].value_counts()
```

```
0
12
```

```
Out[229]: vwc          40507
wug         6515
water board  2933
wua         2535
private operator  1971
parastatal   1768
water authority   904
other         844
company       685
unknown       561
other - school   99
trust         78
Name: management, dtype: int64
```

```
In [230]: print(df['management_group'].isna().sum())
print(df['management_group'].nunique())
df['management_group'].value_counts()
```

```
0
5
```

```
Out[230]: user-group    52490
commercial    3638
parastatal    1768
other         943
unknown       561
Name: management_group, dtype: int64
```

the columns 'management' & 'scheme_management' has almost similar categorical values. So it is better to drop one of them. as can be seen above, 'management' has zero missing values whereas 'scheme_management' has 3877 missing values. In view of this we will drop the column 'scheme_management'

```
In [231]: df.drop(columns=['scheme_management'], inplace=True)
```

```
In [232]: df.shape
```

```
Out[232]: (59400, 34)
```

```
In [233]: df.groupby(['management_group', 'management']).size()
```

```
Out[233]: management_group management
commercial      company           685
                private operator  1971
                trust             78
                water authority   904
other           other            844
                other - school    99
parastatal      parastatal       1768
unknown        unknown          561
user-group      vwc              40507
                water board      2933
                wua              2535
                wug              6515

dtype: int64
```

From above analysis it is understood that the column 'management_group' contains group of categories present in the column 'management'. Hence we will drop the column 'management_group' which has less information as compared to 'management'

```
In [234]: df.drop(columns='management_group', inplace=True)
```

```
In [235]: df.shape
```

```
Out[235]: (59400, 33)
```

Further, the column 'scheme_name' has 28166 missing values and 2696 categories. We will drop the same for now.

```
In [236]: df.drop(columns='scheme_name', inplace=True)
```

```
In [237]: df.shape
```

```
Out[237]: (59400, 32)
```

3.5.2.16 Column "permit"

```
In [238]: df['permit'].value_counts()
```

```
Out[238]: True      38852
          False    17492
          Name: permit, dtype: int64
```

```
In [239]: df['permit'].isna().sum()
```

```
Out[239]: 3056
```

There are 3334 missing values, we will replace those with most occurring values i.e. 'True'

```
In [240]: df['permit'].fillna(value=True, inplace=True)
```

```
In [241]: df['permit'].value_counts()
```

```
Out[241]: True      41908  
False    17492  
Name: permit, dtype: int64
```

3.5.2.17 Column "construction_year"

```
In [242]: print(df['construction_year'].isna().sum())  
print(df['construction_year'].nunique())
```

```
0  
55
```

```
In [243]: df['construction_year'].value_counts()
```

```
Out[243]: 0      20709  
2010     2645  
2008     2613  
2009     2533  
2000     2091  
2007     1587  
2006     1471  
2003     1286  
2011     1256  
2004     1123  
2012     1084  
2002     1075  
1978     1037  
1995     1014  
2005     1011  
1999      979  
1998      966  
1990      954  
1985      945
```


We can make use of this column in conjunction with the column 'date_recorded'. We can get time difference between the year which the waterpoint was constructed in and the year which the data was recorded in. This time difference is nothing but the time which the waterpoint has been operational for. This way we can create a new feature called 'Operational_years'

first we will replace the missing values present in 'construction_year'

```
In [244]: df.loc[df['construction_year']!=0].describe() #to know the statistic of feature without zero values
```

Out[244]:

	id	amount_tsh	gps_height	longitude	latitude	district_code	population	construction_year
count	38691.000000	38691.000000	38691.000000	38691.000000	38691.000000	38691.000000	38691.000000	38691.000000
mean	37083.008736	466.457534	1002.367760	35.983262	-6.235372	5.969786	279.585470	1996.814686
std	21420.922010	3541.036030	618.078669	2.558709	2.761317	10.700673	549.961837	12.472045
min	1.000000	0.000000	-63.000000	29.607122	-11.649440	1.000000	1.000000	1960.000000
25%	18489.500000	0.000000	372.000000	34.676719	-8.755274	2.000000	40.000000	1987.000000
50%	37078.000000	0.000000	1154.000000	36.648187	-6.064216	3.000000	150.000000	2000.000000
75%	55514.500000	200.000000	1488.000000	37.803940	-3.650661	5.000000	305.000000	2008.000000
max	74247.000000	350000.000000	2770.000000	40.345193	-1.042375	63.000000	30500.000000	2013.000000

```
In [245]: df['construction_year'].replace(to_replace = 0, value = 1996, inplace=True)
#replacing the missing values in construction_year column with mean value i.e. 1996
```

```
In [246]: df.describe()
```

Out[246]:

	id	amount_tsh	gps_height	longitude	latitude	district_code	population	construction_year
count	59400.000000	59400.000000	59400.000000	59400.000000	5.940000e+04	59400.000000	59400.000000	59400.000000
mean	37115.131768	317.650385	668.297239	35.116960	-5.706033e+00	5.629747	281.087167	1996.530657
std	21453.128371	2997.574558	693.116350	2.573963	2.946019e+00	9.633649	451.765813	10.073265
min	0.000000	0.000000	-90.000000	29.607122	-1.164944e+01	0.000000	1.000000	1960.000000
25%	18519.750000	0.000000	0.000000	33.354079	-8.540621e+00	2.000000	100.000000	1996.000000
50%	37061.500000	0.000000	369.000000	34.908743	-5.021597e+00	3.000000	281.087167	1996.000000
75%	55656.500000	20.000000	1319.250000	37.178387	-3.326156e+00	5.000000	281.087167	2004.000000
max	74247.000000	350000.000000	2770.000000	40.345193	-2.000000e-08	80.000000	30500.000000	2013.000000

creating new column 'operational_years'

```
In [247]: print(df.date_recorded.head(5))  
print(df.construction_year.head(5))
```

```
0    2011-03-14  
1    2013-03-06  
2    2013-02-25  
3    2013-01-28  
4    2011-07-13  
Name: date_recorded, dtype: object  
0    1999  
1    2010  
2    2009  
3    1986  
4    1996  
Name: construction_year, dtype: int64
```

```
In [248]: df['date_recorded'] = pd.to_datetime(df['date_recorded']) #converting dates to 'datetime' datatype
```

```
In [249]: df['date_recorded'].head()
```

```
Out[249]: 0    2011-03-14  
1    2013-03-06  
2    2013-02-25  
3    2013-01-28  
4    2011-07-13  
Name: date_recorded, dtype: datetime64[ns]
```

```
In [250]: df.date_recorded.dt.year.head(5)
```

```
Out[250]: 0    2011  
1    2013  
2    2013  
3    2013  
4    2011  
Name: date_recorded, dtype: int64
```

```
In [251]: df['operational_year'] = df.date_recorded.dt.year - df.construction_year  
df.operational_year.head(5)
```

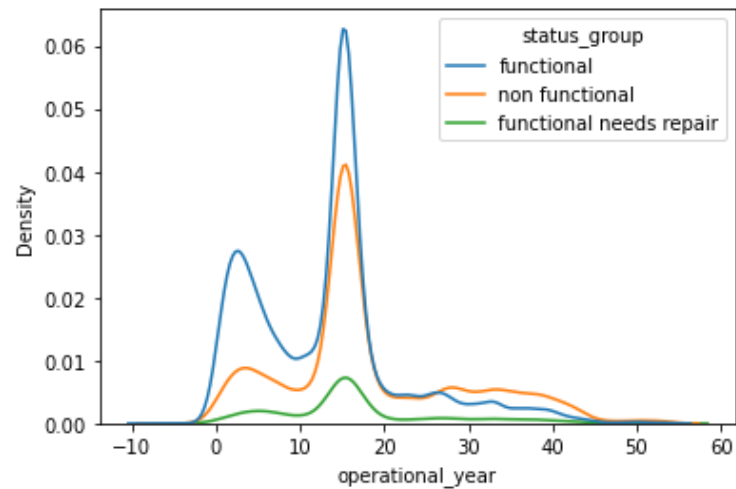
```
Out[251]: 0      12  
          1       3  
          2       4  
          3      27  
          4      15  
          Name: operational_year, dtype: int64
```

```
In [252]: df['operational_year'].value_counts()
```

```
Out[252]: 15      14336  
          16       5968  
          17       2846  
           3       2740  
           1       2303  
           2       2129  
           5       1980  
           4       1890  
          13       1869  
           7       1404  
           6       1382  
          11       1352  
           8       1173  
          14       1160  
          33       1120  
          23        905  
          10        868  
           9        814  
          19        766  
          27        760
```

```
In [253]: sns.kdeplot(data=df, x='operational_year', hue="status_group", gridsize=200)
```

```
Out[253]: <AxesSubplot:xlabel='operational_year', ylabel='Density'>
```



```
In [254]: # plt.figure(figsize=(26,15))  
# ax = sns.countplot(data=df, x='operational_year', hue="status_group")
```

There are some anomalies in data as some of the values in 'operational_year' are negative and operational years can't be negative. We will replace the negative values with minimum value.

```
In [255]: df.loc[~df['operational_year']<0].describe()
```

```
Out[255]:
```

	id	amount_tsh	gps_height	longitude	latitude	district_code	population	construction_year	operational_year
count	59391.000000	59391.000000	59391.000000	59391.000000	5.939100e+04	59391.000000	59391.000000	59391.000000	59391.000000
mean	37116.883753	317.687240	668.297688	35.116650	-5.705857e+00	5.628597	281.086305	1996.528919	15.393949
std	21452.886206	2997.799583	693.118358	2.573822	2.946049e+00	9.632382	451.785541	10.073017	10.089123
min	0.000000	0.000000	-90.000000	29.607122	-1.164944e+01	0.000000	1.000000	1960.000000	0.000000
25%	18522.500000	0.000000	0.000000	33.353967	-8.540784e+00	2.000000	100.000000	1996.000000	8.000000
50%	37063.000000	0.000000	369.000000	34.908362	-5.021241e+00	3.000000	281.087167	1996.000000	15.000000
75%	55661.000000	20.000000	1319.000000	37.177970	-3.326129e+00	5.000000	281.087167	2004.000000	17.000000
max	74247.000000	350000.000000	2770.000000	40.345193	-2.000000e-08	80.000000	30500.000000	2013.000000	53.000000

we will replace the negative values in 'operational_year' with the minimum value which is zero

```
In [256]: df.loc[df['operational_year']<0, 'operational_year'] = 0
```


```
In [257]: df['operational_year'].value_counts().head(20)
```

```
Out[257]: 15    14336
          16     5968
          17     2846
           3     2740
           1     2303
           2     2129
           5     1980
           4     1890
          13     1869
           7     1404
           6     1382
          11     1352
           8     1173
          14     1160
          33     1120
          23      905
          10      868
           9      814
          19      766
```

```
In [258]: # f = lambda x: 0 if x<0 else 1
# df['my_column'] = df['my_column'].map(f) #above can be done using map fuction
```

```
In [259]: sns.kdeplot(data=df, x='operational_year', hue="status_group", gridsize=200)
```

```
Out[259]: <AxesSubplot:xlabel='operational_year', ylabel='Density'>
```



From above plot, we can see that the negative values have disappeared. We will now drop columns 'construction_year' & 'date_recorded'

```
In [260]: df.drop(columns=["construction_year", "date_recorded"], inplace=True)
```

```
In [261]: df.shape
```

```
Out[261]: (59400, 31)
```

```
In [262]: df.columns
```

```
Out[262]: Index(['id', 'amount_tsh', 'funder', 'gps_height', 'installer', 'longitude',
               'latitude', 'basin', 'region', 'district_code', 'lga', 'population',
               'public_meeting', 'permit', 'extraction_type', 'extraction_type_group',
               'extraction_type_class', 'management', 'payment', 'payment_type',
               'water_quality', 'quality_group', 'quantity', 'quantity_group',
               'source', 'source_type', 'source_class', 'waterpoint_type',
               'waterpoint_type_group', 'status_group', 'operational_year'],
              dtype='object')
```

3.5.2.18 Column "extraction_type", "extraction_type_group", "extraction_type_class"


```
In [263]: print(df['extraction_type'].isna().sum())
print(df['extraction_type'].nunique())
df['extraction_type'].value_counts()
```

0

18

```
Out[263]: gravity                26780
nira/tanira                    8154
other                         6430
submersible                   4764
swn 80                        3670
mono                          2865
india mark ii                 2400
afridev                      1770
ksb                           1415
other - rope pump             451
other - swan 81               229
windmill                     117
india mark iii                98
cemo                          90
other - play pump             85
walimi                       48
climax                       32
other - mkulima/shinyanga     2
Name: extraction_type, dtype: int64
```

```
In [264]: plt.figure(figsize=(15,7))
ax = sns.countplot(x='extraction_type', hue="status_group", data=df)
ax.tick_params(axis='x', rotation=45)
# plt.xticks(rotation=90)
```




```
In [265]: print(df['extraction_type_group'].isna().sum())
print(df['extraction_type_group'].nunique())
df['extraction_type_group'].value_counts()
```

```
0
13
```

```
Out[265]: gravity          26780
nira/tanira          8154
other                6430
submersible         6179
swn 80              3670
mono                2865
india mark ii       2400
afridev            1770
rope pump           451
other handpump       364
other motorpump      122
wind-powered        117
india mark iii       98
Name: extraction_type_group, dtype: int64
```

```
In [266]: print(df['extraction_type_class'].isna().sum())
print(df['extraction_type_class'].nunique())
df['extraction_type_class'].value_counts()
```

```
0
7
```

```
Out[266]: gravity          26780
handpump              16456
other                6430
submersible         6179
motorpump           2987
rope pump            451
wind-powered        117
Name: extraction_type_class, dtype: int64
```

```
In [267]: df.groupby(['extraction_type_class', 'extraction_type']).size()
```

```
Out[267]: extraction_type_class  extraction_type
gravity                          gravity          26780
handpump                        afridev           1770
                                india mark ii      2400
                                india mark iii       98
                                nira/tanira        8154
                                other - mkulima/shinyanga  2
                                other - play pump     85
                                other - swan 81       229
                                swan 80            3670
                                walimi              48
motorpump                       cemo              90
                                climax              32
                                mono              2865
other                           other            6430
rope pump                       other - rope pump   451
submersible                     ksb              1415
                                submersible        4764
wind-powered                    windmill          117
dtype: int64
```

```
In [268]: df.groupby(['extraction_type_class', 'extraction_type_group']).size()
```

```
Out[268]: extraction_type_class  extraction_type_group
gravity                          gravity          26780
handpump                        afridev           1770
                                india mark ii      2400
                                india mark iii       98
                                nira/tanira        8154
                                other handpump       364
                                swan 80            3670
motorpump                       mono            2865
                                other motorpump      122
other                           other            6430
rope pump                       rope pump         451
submersible                     submersible     6179
wind-powered                    wind-powered     117
dtype: int64
```

"extraction_type", "extraction_type_group" & "extraction_type_class" provide the details as regard the type of pumping system being used. Here, 'extraction_type_class' is sort of group version of "extraction_type" & "extraction_type_group". Here we will chose to keep "extraction_type_group" and remove other two. The reason we are removing "extraction_type" is because it has few categories which have very less no. of values.

```
df.drop(columns=["extraction_type", "extraction_type_class"], inplace=True)
```

```
df.shape
```


(59400, 29)

3.5.2.19 Column "payment", "payment_type"

```
print(df['payment'].isna().sum())
print(df['payment'].nunique())
df['payment'].value_counts()
```

```
never pay          25348
pay per bucket     8985
pay monthly        8300
unknown            8157
pay when scheme fails 3914
pay annually       3642
other              1054
Name: payment, dtype: int64
```

```
In [272]: plt.figure(figsize=(10,5))
ax = sns.countplot(x='payment', hue="status_group", data=df)
ax.tick_params(axis='x', rotation=45)
```



```
In [273]: print(df['payment_type'].isna().sum())
print(df['payment_type'].nunique())
df['payment_type'].value_counts()
```

```
0
7
```

```
Out[273]: never pay      25348
per bucket    8985
monthly       8300
unknown       8157
on failure    3914
annually      3642
other         1054
Name: payment_type, dtype: int64
```

Here both features are providing similar information. We will remove one of them.

```
In [274]: df.drop(columns='payment_type', inplace=True)
```

```
In [275]: df.shape
```

```
Out[275]: (59400, 28)
```

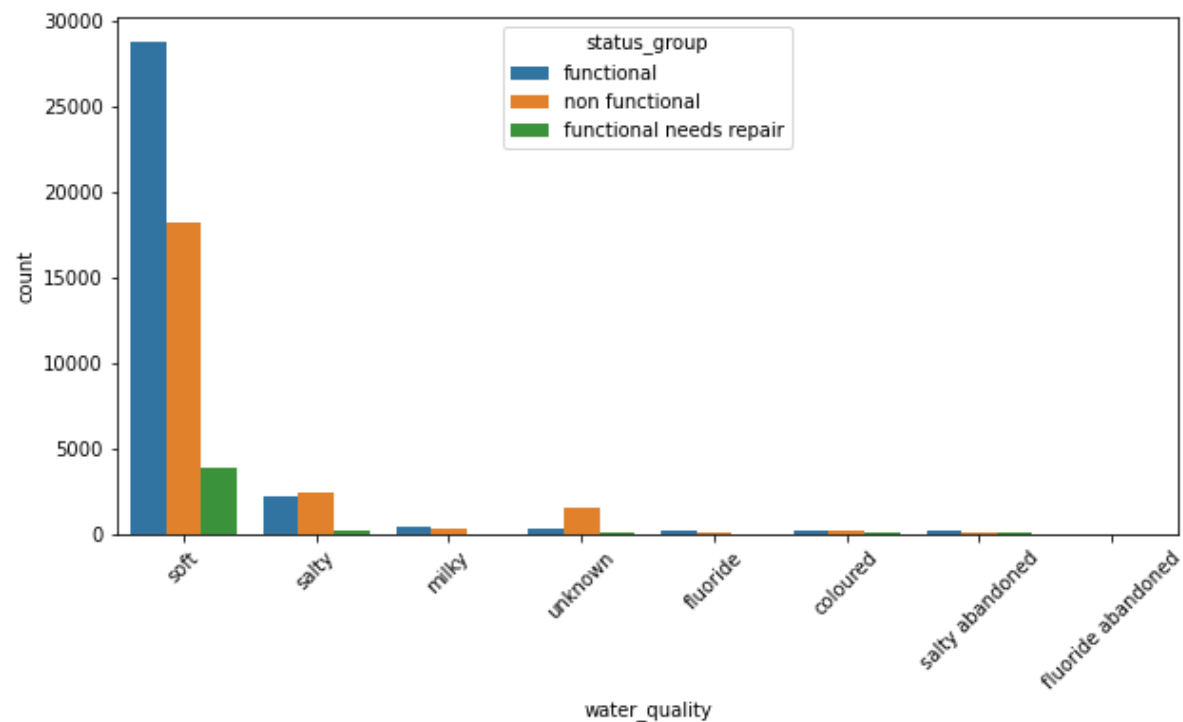
3.5.2.20 Column "water_quality", "quality_group", "quantity", "quantity_group"

```
In [276]: print(df['water_quality'].isna().sum())
print(df['water_quality'].nunique())
df['water_quality'].value_counts()
```

```
0
8
```

```
Out[276]: soft          50818
salty          4856
unknown        1876
milky           804
coloured        490
salty abandoned  339
fluoride        200
fluoride abandoned  17
Name: water_quality, dtype: int64
```

```
In [277]: plt.figure(figsize=(10,5))
ax = sns.countplot(x='water_quality', hue="status_group", data=df)
ax.tick_params(axis='x', rotation=45)
```



```
In [278]: print(df['quality_group'].isna().sum())
print(df['quality_group'].nunique())
df['quality_group'].value_counts()
```

```
0
6
```

```
Out[278]: good      50818
salty      5195
unknown    1876
milky      804
colored    490
fluoride   217
Name: quality_group, dtype: int64
```


The columns are similar. The column "water_quality" tend to give more informayion than that of "quality_group" so we will drop "quality_group"

```
In [279]: print(df['quantity'].isna().sum())  
print(df['quantity'].nunique())  
df['quantity'].value_counts()
```

```
0  
5
```

```
Out[279]: enough          33186  
insufficient      15129  
dry                6246  
seasonal          4050  
unknown           789  
Name: quantity, dtype: int64
```

```
In [280]: plt.figure(figsize=(10,5))  
ax = sns.countplot(x='quantity', hue="status_group", data=df)  
ax.tick_params(axis='x', rotation=45)
```



```
In [281]: print(df['quantity_group'].isna().sum())
print(df['quantity_group'].nunique())
df['quantity_group'].value_counts()
```

```
0
5
```

```
Out[281]: enough          33186
insufficient    15129
dry             6246
seasonal        4050
unknown         789
Name: quantity_group, dtype: int64
```

quantity_group & quantity are similar. We will drop one of them

```
In [282]: df.drop(columns=["quality_group", "quantity_group"], inplace=True)
```

```
In [283]: df.shape
```

```
Out[283]: (59400, 26)
```

3.5.2.21 Column "source", "source_type", "source_class"

```
In [284]: print(df['source'].isna().sum())
print(df['source'].nunique())
df['source'].value_counts()
```

```
0
10
```

```
Out[284]: spring          17021
shallow well            16824
machine dbh             11075
river                   9612
rainwater harvesting     2295
hand dtw                 874
lake                    765
dam                     656
other                   212
unknown                  66
Name: source, dtype: int64
```



```
In [285]: print(df['source_type'].isna().sum())
print(df['source_type'].nunique())
df['source_type'].value_counts()
```

```
0
7
```

```
Out[285]: spring          17021
shallow well          16824
borehole              11949
river/lake            10377
rainwater harvesting   2295
dam                   656
other                  278
Name: source_type, dtype: int64
```

```
In [286]: print(df['source_class'].isna().sum())
print(df['source_class'].nunique())
df['source_class'].value_counts()
```


```
0
3
```

```
Out[286]: groundwater    45794
surface                 13328
unknown                  278
Name: source_class, dtype: int64
```

```
In [287]: df.groupby(['source_class', 'source']).size()
```

```
Out[287]: source_class  source
groundwater  hand dtw          874
              machine dbh      11075
              shallow well      16824
              spring           17021
surface      dam               656
              lake              765
              rainwater harvesting 2295
              river             9612
unknown      other             212
              unknown           66
dtype: int64
```

```
In [288]: plt.figure(figsize=(10,5))
ax = sns.countplot(x='source', hue='status_group', data=df)
ax.tick_params(axis='x', rotation=45)
```



Among above 3 columns we will choose to keep 'source' as it has more information. We will drop rest two columns.

```
In [289]: df.drop(columns=["source_type", "source_class"], inplace=True)
```

```
In [290]: df.shape
```

```
Out[290]: (59400, 24)
```

3.5.2.22 Column "waterpoint_type", "waterpoint_type_group"

```
In [291]: print(df['waterpoint_type'].isna().sum())  
print(df['waterpoint_type'].nunique())  
df['waterpoint_type'].value_counts()
```

```
0  
7
```


```
Out[291]: communal standpipe      28522  
hand pump      17488  
other      6380  
communal standpipe multiple      6103  
improved spring      784  
cattle trough      116  
dam      7  
Name: waterpoint_type, dtype: int64
```

```
In [292]: print(df['waterpoint_type_group'].isna().sum())  
print(df['waterpoint_type_group'].nunique())  
df['waterpoint_type_group'].value_counts()
```

```
0  
6
```

```
Out[292]: communal standpipe      34625  
hand pump      17488  
other      6380  
improved spring      784  
cattle trough      116  
dam      7  
Name: waterpoint_type_group, dtype: int64
```

```
In [293]: plt.figure(figsize=(10,5))
ax = sns.countplot(x='waterpoint_type', hue="status_group", data=df)
ax.tick_params(axis='x', rotation=45)
```



here we will drop the column "waterpoint_type_group" as it has less information

```
In [294]: df.drop(columns=["waterpoint_type_group"], inplace=True)
```

```
In [295]: df.columns
```

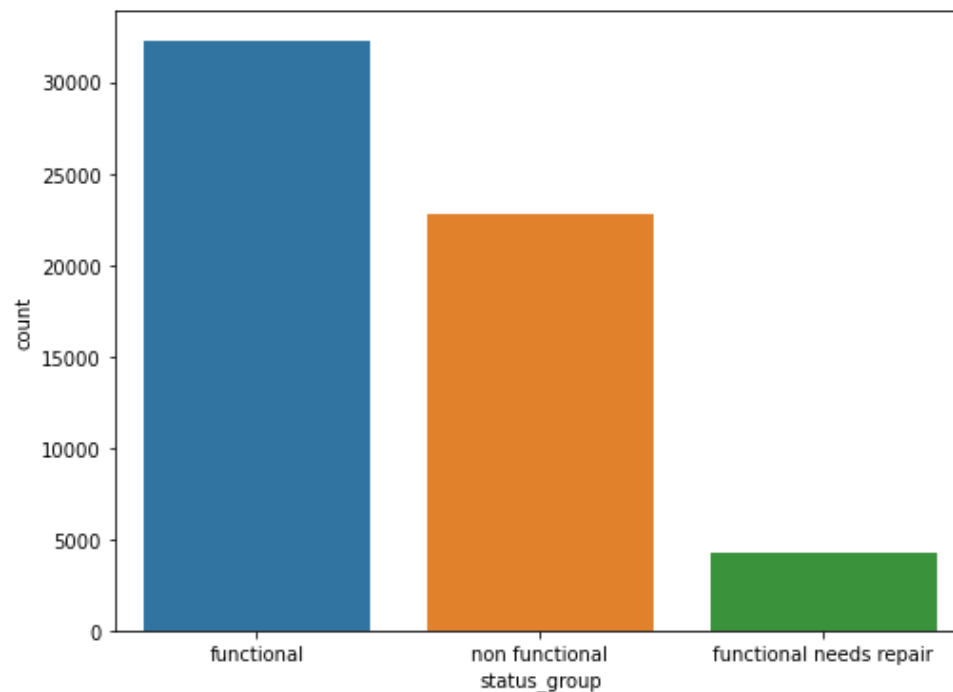
```
Out[295]: Index(['id', 'amount_tsh', 'funder', 'gps_height', 'installer', 'longitude',  
               'latitude', 'basin', 'region', 'district_code', 'lga', 'population',  
               'public_meeting', 'permit', 'extraction_type_group', 'management',  
               'payment', 'water_quality', 'quantity', 'source', 'waterpoint_type',  
               'status_group', 'operational_year'],  
              dtype='object')
```

```
In [296]: df.shape
```

```
Out[296]: (59400, 23)
```

3.5.2.23 Column "status_group"

```
In [297]: plt.figure(figsize=(8,6))  
ax = sns.countplot(x="status_group", data=df)
```



From above plot of class labels i.e. "status_group", it is understood that the data is highly imbalanced.

In [298]: df.head(5)

Out[298]:

	id	amount_tsh	funder	gps_height	installer	longitude	latitude	basin	region	district_code	lga	population	public_meeting	permit	e
0	69572	6000.0	other	1390	other	34.938093	-9.856322	Lake Nyasa	Iringa	5	Ludewa	109.000000	True	False	
1	8776	0.0	other	1399	other	34.698766	-2.147466	Lake Victoria	Mara	2	Serengeti	280.000000	True	True	
2	34310	25.0	other	686	World Vision	37.460664	-3.821329	Pangani	Manyara	4	Simanjiro	250.000000	True	True	
3	67743	0.0	Unicef	263	UNICEF	38.486161	-11.155298	Ruvuma / Southern Coast	Mtwara	63	Nanyumbu	58.000000	True	True	
4	19728	0.0	other	0	other	31.130847	-1.825359	Lake Victoria	Kagera	1	Karagwe	281.087167	True	True	



In [299]: df.shape

Out[299]: (59400, 23)

3.6 exporting cleaned data to CSV

In [300]: df.to_csv('clean_df.csv') *#exporting clean data to csv*

In []: